

**Bottom-Up and Top-Down Attention for Image Captioning and Visual Question Answering:** trained with Adam (learning rate = 0.001) and L2 regularization on all the weights except for the bias with gamma =  $5 \cdot 10^{-6}$

